

Some adversarial attempts on MNIST

1. Abstract

This project

1. Trained a model (will be referred as original model later) to predict the handwritten digits (MNIST) dataset
2. Collect the flip of points on image, to make the model recognize the image to a wrong number
3. Explain that the greedy approach for collecting flips is reasonable
4. Train the adversarial model based on the collected flips
5. Use the adversarial model to create examples to mislead original model and validate the attack success rate
6. Train a distillation model from original model
7. Compare the recognition on adversarial examples generated from adversarial model for original model and distillation model

2. Experiments

1. The model for MNIST comes directly from Keras. The file locates at `model.py`

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 64)	0
flatten (Flatten)	(None, 1600)	0
dropout (Dropout)	(None, 1600)	0
dense (Dense)	(None, 10)	16010

Total params: 34,826

Trainable params: 34,826

Non-trainable params: 0

2. Flip of points

There is explanation why to use flips (rather than changing greyscale in certain ways) in progress report 2(I missed the due date and when I finished it I found only 2 people submitted, so I am just attaching the report along with this final report). After around 7 days running at server (Intel(R) Xeon(R) Platinum 8260 CPU @ 2.40GHz, not any fancy performance CPU, but not bad) collected 7000 sets of points to attack the original model.

The file used in this part is `run.py`

3. Greedy explanation

Assume if we have the points, we want to prove that the greedy approach (find each flip of points that lower the original confidence most) will yield the least amount of flips to mislead the original model. We design a small experiment to explain that greedy approach will produce **reasonable small amount** of flips to mislead the original model.

Assume we have points, we pin 2 random points from the sequential lowest to the collection, then use the same approach to find the flips to lower the original model confidence. Then we compare the random pinned points to the original sequential points, to see the difference.

The result is that, if we pin 2 random points of sequential points and find the flips, the same points versus the different points is 332-140 (for small amount of tests as finding these points is very time-consuming). The points of same amount of points and less versus larger is 68-7. That is, with our greedy approach, it will produce reasonably least amount of points to lower the original model confidence.

The source code for testing locates at `greedy_check.py` and `greedy_compare.py`

4. Train adversarial model

From above we have collected the points to attack the original model, we want to train an adversarial model on the original model. An observation from the points we collected is that on average, we need 6.7 points to make original model recognize the image as wrong number.

I tried train the model to predict x positions (tried 5 and 7) that will lower the original model confidence, that is the model input shape to output shape $28*28*1 \rightarrow x*2*1$ However, after many attempts, I can't managed to train the model with a reasonable accuracy, most of the time the models will converge the accuracy at 0.114, I am not sure why.

Then, I tried to train a model that have the points marked as 1 laying on empty $28*28$ zero array, then use similar layers in original model, to create the adversarial model.


Model: "sequential_21"

Layer (type)	Output Shape	Param #
conv2d_31 (Conv2D)	(None, 25, 25, 32)	544
max_pooling2d_40 (MaxPooling2D)	(None, 12, 12, 32)	0
flatten_22 (Flatten)	(None, 4608)	0
dropout_20 (Dropout)	(None, 4608)	0
dense_19 (Dense)	(None, 784)	3613456
reshape_14 (Reshape)	(None, 28, 28, 1)	0
Total params: 3,614,000		
Trainable params: 3,614,000		
Non-trainable params: 0		


5. To create adversarial examples

The output of such model on a 28x28 array is a 28x28 array, and the positions range from 0-1, so we need a factor that multiplies the position. For example, assume we have 0.13 at position (2,5), we set the multiplier to be 20, then it becomes 2.6 at point (2,5), we limit each point at max 1, so point (2,5) becomes 1, other points if less than 1 we change that point to 0 (by converting each point to type int). Test run the images and see how big `multiple_factor` will change the recognition from original model, the higher this factor is, our model will need more points to mislead the original model.

After test on around 2000 images, I found that the average for such multiplier factor is 276, and most of the image will share a similar shape for attack. If we feed the output from the adversarial model to heatmap, we can see that at which position attack will more likely happen (at least the model thinks so) For example, if we have a image that looks like this



Original 5 Now 9



Original 6 Now 8

6. Train distillation model

Based on the paper Distillation as a Defense to Adversarial Perturbations against Deep Neural Networks

I first processed the `y_train` predicted by original model (10*1 array such as [0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 0.1]), then use it as `y` to train the new model

Model: "sequential_38"

Layer (type)	Output Shape	Param #
conv2d_61 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_70 (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_62 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_71 (MaxPooling2D)	(None, 5, 5, 64)	0
flatten_41 (Flatten)	(None, 1600)	0


dropout_36 (Dropout)	(None, 1600)	0
dense_39 (Dense)	(None, 10)	16010

```
=====
Total params: 34,826
Trainable params: 34,826
Non-trainable params: 0
```


7. Attack the distillation model

Recall the `multiply_factor` that controls the points added to the attacked image, this number reflects how **hard** to attack the model. As we have trained the distillation model, we want to validate that whether the proposed distillation model is more robust against attacks, so we randomly check what is the `multiply_factor` to change the prediction of original model and distillation model, compare the `multiply_factor` we can see the robustness of these 2 models.

For example, the original image and the heatmap for attack points



The minimum `multiply_factor` to attack the original model is shown on left, and the minimum `multiply_factor` to attack the distillation model is shown on the right side



We can see it need few more points to attack the distillation model

Then, we run tests on these 2 models and check the average `multiply_factor` for original model and distillation model.

And check 600 random images out of the train images, that the average factor to mislead the original model is 23, while the average for distillation

model is 27, which shows the distillation model has better robustness than the original model

References: [1] Papernot, Nicolas, et al. "Distillation as a defense to adversarial perturbations against deep neural networks." 2016 IEEE symposium on security and privacy (SP). IEEE, 2016.

[2] Athalye, Anish, Nicholas Carlini, and David Wagner. "Obfuscated gradients give a false sense of security: Circumventing defenses to adversarial examples." International conference on machine learning. PMLR, 2018.

[3] MNIST Adversarial Examples Challenge https://github.com/MadryLab/mnist_challenge

[4] mnist-adversary <https://github.com/dguliani/mnist-adversary>

[5] Simple MNIST convnet https://keras.io/examples/vision/mnist_convnet/

[6] Help from friend Sida Zhu